

# T1 Graph Concepts

**Szilard Barany**  
Principal Sales Engineer

---

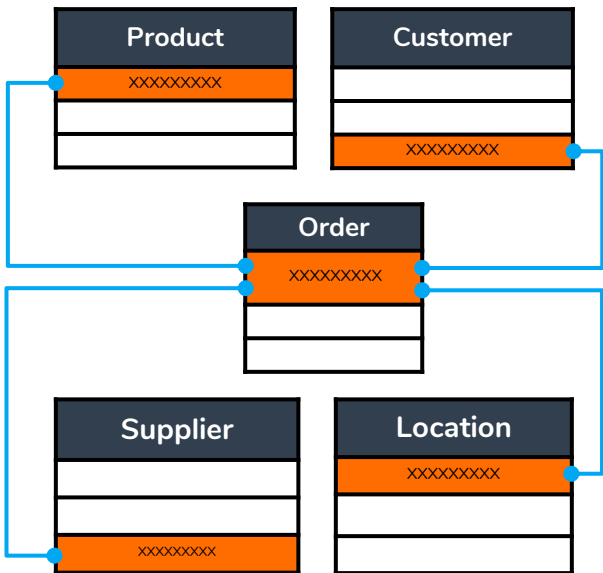


# Agenda

- Evolution of databases
- Evolution of graph databases
- Graph concepts
- Data science capabilities of a graph analytical platform


# The Evolution of Databases

## Relational Database



- Rigid schema
- High performance for transactions
- Poor performance for deep analytics

## Key-Value Database





















- Highly fluid schema/no schema
- High performance for simple transactions
- Poor performance deep analytics

## Graph Database



- Flexible schema
- High performance for complex transactions
- High performance for deep analytics

# The Evolution of Graph Databases

	<b>Graph 1.0</b> Single server, non-parallel	<b>Graph 2.0</b> NoSQL base for storage scale	<b>Graph 3.0</b> Native, Parallel
<b>Native Graph Storage</b>		 Key-value or column store	
<b>Parallel Loading</b>	 Days to load terabytes	 Days to load terabytes	 Hours to load terabytes
<b>Parallel Multi-Hop Analytics</b>	 Times out after 2 hops	 Runs out of time/memory after 2 hops	 Sub-second across 10+ hops
<b>Parallel Updates (in real-time)</b>	 Batch updates	 Batch updates	 Mutable/Transactional
<b>Scale Out for Speed and Size</b>	 Requires manual sharding of data and queries	 Sharded graphs, handles 1-2 hop queries	 Auto distributed data & analytics, for TBs data & 2B+ deep queries/day
<b>Privacy for Sensitive Data</b>			 MultiGraph

# TigerGraph: 3rd Generation Graph Database

## Real-time Performance


Sub-second response for queries touching tens of millions of entities/relationships.

## Transactional (Mutable) Graph

- Hybrid Transactional / Analytical Processing (HTAP): OLTP and OLAP.
- Hundreds of thousands of updates per second, billions of transactions per day.

## Scalability for Massive Datasets

100 B+ entities, 1 T+ relationships.



## Deep Link Multi-Hop Analytics

- Queries traverse 10+ hops deep into the graph performing complex calculations.

## Ease of Development & Deployment

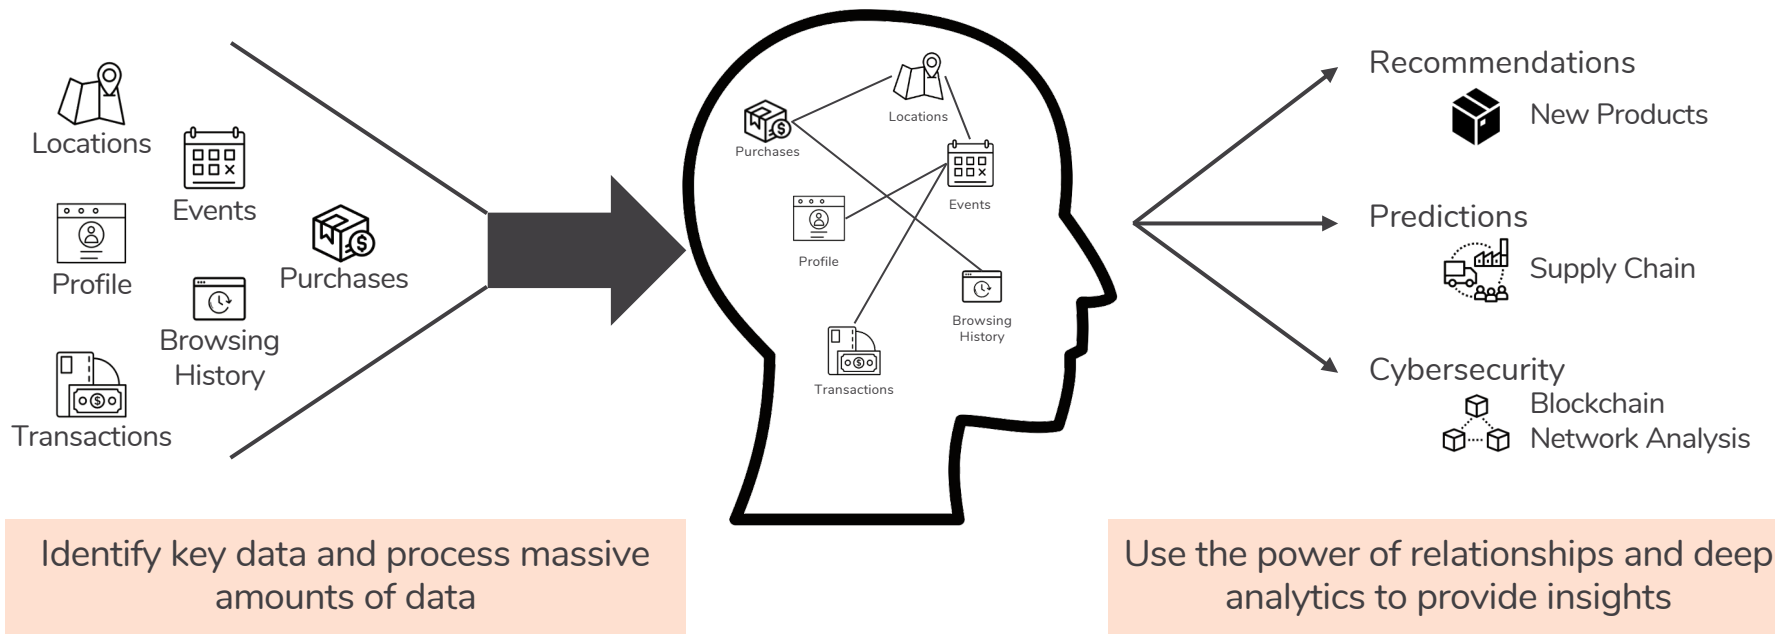
- GraphStudio - visual SDK.
- GSQL - Intuitive, Turing complete graph query language for developing complex analytics in days.
- User extensible graph algorithms library.

## Enterprise Grade Security

- Encryption support.
- Control access to sensitive data based on user role, department or organization with MultiGraph.

# Graph is How We Think

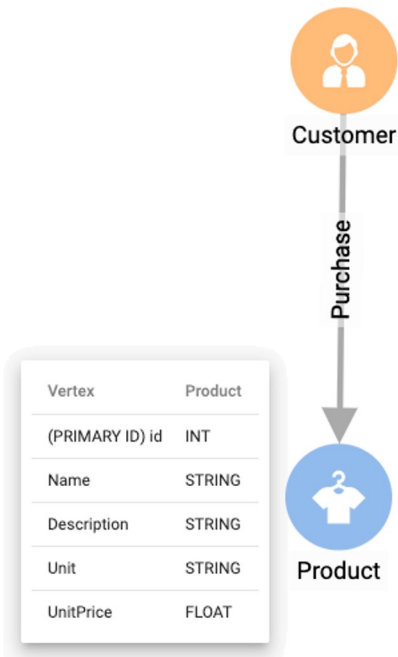
Graph is a natural model for interconnected data.  
It is an organic way of modeling data for a variety of  
relationships and transactions.



# What is a graph?

A collection of

- Vertices (vertexes, nodes) and
  - Edges (links, relationships) between vertices
- 
- A graph is an object consisting of two sets called vertex set and its edge set.
  - The vertex set is a finite nonempty set.
  - The edge set may be empty, but otherwise its elements are two-element subsets of the vertex set.
  - The elements of the vertex set of a graph are called vertices and the elements of the edge set are called edges.



# What is a vertex? What is an edge?

- Vertices represent real life entities
  - People, products, parts, IP addresses, genre, location, date, ...
  - Visualised as circle/disc
  - Can be considered similar to a relational database table, but ...
  - Can be called nodes, but could that be unambiguous
- Edges represent relationships between real life entities
  - Parenthood, friendship, ownership, hierarchy, action, ...
  - Visualised as line or arrow
  - Can be considered similar to a join in a relational database, but ...
  - Sometimes a relationship can be better represented as a vertex




# Vertices and Edges

- Vertices and edges
  - Have type
  - Can have attributes
- Vertices have (primary) IDs
  - Unique value within a vertex type
  - If you insert a vertex with the same ID, it will overwrite the previous one
  - Similar to primary key in relational databases
- Edges do not have IDs
  - They have source and target vertices (vertex IDs)
  - Their combination identifies the edge uniquely


# Seven Key Data Science Capabilities

## 1 Deep Link Analysis




For a set of entities (e.g., customers, accounts, citizens, doctors), show all links or connections

## 2 Subgraph or Relationship Discovery & Computation




Given a graph (e.g., all payments or claims for a set of customers), find all the relevant subgraphs

## 3 Multi-dimensional Entity & Pattern Matching



Given a graph (e.g., all payments or claims for a set of customers), find similar graphs

## 4 Hub & Community Detection



Find most influential members of a group (customers, doctors, citizens) & detect community around them

## 5 Geospatial graph analysis

Analyze changes in entities & relationships with location data

## 6 Temporal (time-series) graph analysis

Analyze changes in entities & relationships over time

## 7 ML feature generation and explainable AI

Extract graph-based features to feed as training data for ML; power explainable AI

# Community resources

## Developer Community

- TG Community Forum [community.tigergraph.com](https://community.tigergraph.com)
- TG Community Chat [discord.gg/F2c9b9v](https://discord.gg/F2c9b9v)
- Developer Portal <https://dev.tigergraph.com/>
- Reddit [reddit.com/r/tigergraph/](https://reddit.com/r/tigergraph/)
- YouTube [youtube.com/tigergraph](https://youtube.com/tigergraph)
- LinkedIn [linkedin.com/company/tigergraph/](https://linkedin.com/company/tigergraph/)
- Twitter [twitter.com/tigergraphdb](https://twitter.com/tigergraphdb)
- Twitch [twitch.tv/tigergraph](https://twitch.tv/tigergraph)
- GitHub [github.com/tigergraph/ecosys](https://github.com/tigergraph/ecosys)  
[github.com/tigergraph/gsql-graph-algorithms](https://github.com/tigergraph/gsql-graph-algorithms)  
[github.com/TigerGraph-OSS](https://github.com/TigerGraph-OSS)

# Thank you!

Szilard Barany

[szilard.barany@tigergraph.com](mailto:szilard.barany@tigergraph.com)

Discord: Szilard 

